

GradCAM Visualization for CNN Interpretability - Assignment

Objective

The goal of this assignment is to implement and apply **Gradient-weighted Class Activation Mapping (GradCAM)** to visualize and interpret what your trained Convolutional Neural Network has learned for glaucoma detection. You will generate heatmaps showing which regions of retinal images your model focuses on when making predictions.

What is GradCAM?

GradCAM (Gradient-weighted Class Activation Mapping) is a visualization technique that produces visual explanations for decisions from CNN-based models.

The Core Idea

Problem: CNNs are black boxes - we can't see what they're "looking at" when making predictions.

Solution: GradCAM creates a heatmap overlay on the input image showing:

- **Hot regions (red/yellow):** Areas the model considers IMPORTANT for the prediction
- **Cold regions (blue/green):** Areas the model considers LESS important

How does it work? (High-level)

1. Forward pass through the network
2. Get the prediction for a specific class
3. Compute gradients of that class score with respect to feature maps
4. Weight the feature maps by these gradients
5. Create a heatmap showing important regions

You don't need to understand all the math details - the key is understanding what it shows you!

Why is GradCAM Important?

1. Model Validation

- Verify the model looks at clinically relevant regions
- Check if it's using shortcuts or artifacts
- Ensure it focuses on the optic disc for glaucoma

2. Trust and Transparency

- Show doctors what the AI is "seeing"
- Build confidence in predictions

- Meet explainability requirements

3. Error Analysis

- Understand why the model makes mistakes
- Identify biases in the model
- Find problematic images or patterns

4. Clinical Insights

- Compare model attention with expert knowledge
- Discover if CNN found novel patterns
- Validate medical relevance of predictions

5. Debugging

- Check if model learned correct features
- Identify if it's overfitting to spurious correlations
- Guide model improvements

Medical Context: Glaucoma Detection

What should the model look at?

For glaucoma diagnosis, ophthalmologists examine:

1. Optic Disc (Optic Nerve Head)

- The circular/oval structure where the optic nerve enters the eye
- Primary area of interest for glaucoma

2. Optic Cup

- The bright central area within the disc
- Enlarged cup is a key glaucoma indicator

3. Neuroretinal Rim

- The orange/pink tissue between cup and disc edge
- Thinning indicates glaucoma progression

4. Cup-to-Disc Ratio (CDR)

- Ratio of cup diameter to disc diameter
- CDR > 0.5-0.6 suggests glaucoma

5. RNFL (Retinal Nerve Fiber Layer)

- Nerve fibers radiating from optic disc
- Thinning or defects indicate damage

What SHOULDN'T the model focus on?

- Blood vessels away from the disc
 - Image edges or artifacts
 - Imaging equipment markers
 - Peripheral retina (for glaucoma)
-

The Task

What You'll Work With

1. Trained CNN Model

- Your best model from the competition
- Already trained on glaucoma detection

2. Test Images

- Retinal fundus images
- Both normal and glaucoma cases
- Images the model hasn't seen during training

3. Model Predictions

- Class predictions (normal/glaucoma)
- Confidence scores
- Both correct and incorrect predictions

What You'll Generate

For each test image, create visualizations showing:

- Original retinal image
 - GradCAM heatmap
 - Overlay of heatmap on original image
 - Prediction and confidence information
-

Key Concepts

Understanding the Heatmap

Color Interpretation:

- **Red/Yellow:** High activation - model considers this VERY important
- **Green:** Medium activation - some importance
- **Blue/Purple:** Low activation - not important for this prediction

Spatial Meaning: The heatmap shows **WHERE** in the image the model is looking, not necessarily **WHAT** it sees.

Class-specific Visualization

GradCAM is **class-specific**:

- You can generate a heatmap for "normal" class
- You can generate a heatmap for "glaucoma" class
- They will be different (showing what supports each class)

Typically: Generate heatmap for the **predicted class** or the **target class** you want to understand.

Assignment Components

1. Implementation

Implement or adapt a GradCAM visualization system that:

- Loads your trained CNN model
- Processes test images
- Generates GradCAM heatmaps
- Creates visualization overlays
- Saves results systematically

Note: You can use existing libraries (like [pytorch-grad-cam](#)) .

2. Visualization Generation

Generate GradCAM visualizations for:

- **All test images** (or a representative subset)

You can also separate visualizations into:

- **Correctly classified normal** cases
- **Correctly classified glaucoma** cases
- **Misclassified images** (false positives and false negatives)
- **High confidence predictions**
- **Low confidence predictions**

3. Analysis and Interpretation

Analyze the generated visualizations to answer:

Model Behavior Questions

- Where does the model focus for glaucoma predictions?
- Where does the model focus for normal predictions?
- Are the focus regions clinically relevant?
- Does the model look at the optic disc?

Correctness Analysis

- Do correct predictions show appropriate focus?

- What do misclassifications look like?
- Can you identify why the model failed on certain images?

Comparison Analysis

- How do normal vs glaucoma heatmaps differ?
 - Are there patterns across similar images?
 - Does confidence correlate with focus quality?
-

🔧 Technical Considerations

Model Architecture

Different CNN architectures have different internal structures:

Key Concept: You need to identify the **target layer** for GradCAM - typically the last convolutional layer before classification.

Examples:

- **ResNet:** `layer4[-1]` (last block of layer4)
- **DenseNet:** `features[-1]` (last feature layer)
- **EfficientNet:** `features[-1]` (last feature block)
- **VGG:** Last conv layer before flatten

Image Preprocessing

Critical: Use the SAME preprocessing as during training:

- Same normalization values
- Same resize dimensions
- NO data augmentation (use test transforms)

Computational Requirements

- GradCAM requires gradient computation
 - Process one image at a time (or small batches)
 - May take longer than inference alone
 - GPU recommended but not required
-